

In [0]:

```
from keras.utils import np_utils
from keras.datasets import mnist
from keras.initializers import RandomNormal
```

Using TensorFlow backend.

In [0]:

```
import matplotlib.pyplot as plt
import numpy as np
import time
import pandas as pd
import seaborn as sns

import keras
from keras.models import Sequential
from keras.layers import Dense, Activation, Dropout, Flatten
from keras.layers import Conv2D, MaxPooling2D
from keras import backend as K

from keras.layers.normalization import BatchNormalization
from keras.layers import Dropout
```

In [0]:

```
def plt_dynamic(x, vy, ty, ax, colors=['b']):
    ax.plot(x, vy, 'b', label="Validation Loss")
    ax.plot(x, ty, 'r', label="Train Loss")
    plt.legend()
    plt.grid()
    fig.canvas.draw()
```

In [0]:

```
(x_train, y_train), (x_test, y_test) = mnist.load_data()

img_rows = 28
img_cols = 28

if K.image_data_format() == 'channels_first':
    x_train = x_train.reshape(x_train.shape[0], 1, img_rows, img_cols)
    x_test = x_test.reshape(x_test.shape[0], 1, img_rows, img_cols)
    input_shape = (1, img_rows, img_cols)
else:
    x_train = x_train.reshape(x_train.shape[0], img_rows, img_cols, 1)
    x_test = x_test.reshape(x_test.shape[0], img_rows, img_cols, 1)
    input_shape = (img_rows, img_cols, 1)
```

Downloading data from <https://s3.amazonaws.com/img-datasets/mnist.npz>  
11493376/11490434 [=====] - 1s 0us/step

In [0]:

```
x_train = x_train.astype('float32')
x_test = x_test.astype('float32')

x_train = x_train/255
x_test = x_test/255

y_train = np_utils.to_categorical(y_train, 10)
y_test = np_utils.to_categorical(y_test, 10)
```

In [0]:

```
batch_size = 128
```

```

batch_size = 128
num_classes = 10
epochs = 12

```

### 3 Convolutional Layers

In [0]:

```

model = Sequential()

model.add(Conv2D(32, kernel_size=(3, 3),padding='same',activation='relu',input_shape=input_shape))
model.add(Conv2D(64,kernel_size = (5,5),padding='same', activation='relu'))
model.add(MaxPooling2D(pool_size=(2, 2),padding='same'))

model.add(Conv2D(32, kernel_size=(3, 3),padding='same',activation='relu'))
model.add(Conv2D(64,kernel_size = (5,5),padding='same', activation='relu'))
model.add(MaxPooling2D(pool_size=(3,3),padding='same'))

model.add(Conv2D(32, kernel_size=(3, 3),padding='same',activation='relu'))
model.add(Conv2D(64,kernel_size = (5,5), padding='same',activation='relu'))
model.add(MaxPooling2D(pool_size=(4,4),padding='same'))

model.add(Dropout(0.5))

model.add(Flatten())

model.add(Dense(256, activation='relu'))
model.add(BatchNormalization())
model.add(Dropout(0.5))

model.add(Dense(128, activation='relu'))
model.add(BatchNormalization())
model.add(Dropout(0.5))

model.add(Dense(num_classes, activation='softmax'))

model.summary()

```

Layer (type)	Output Shape	Param #
conv2d_36 (Conv2D)	(None, 28, 28, 32)	320
conv2d_37 (Conv2D)	(None, 28, 28, 64)	51264
max_pooling2d_18 (MaxPooling)	(None, 14, 14, 64)	0
conv2d_38 (Conv2D)	(None, 14, 14, 32)	18464
conv2d_39 (Conv2D)	(None, 14, 14, 64)	51264
max_pooling2d_19 (MaxPooling)	(None, 5, 5, 64)	0
conv2d_40 (Conv2D)	(None, 5, 5, 32)	18464
conv2d_41 (Conv2D)	(None, 5, 5, 64)	51264
max_pooling2d_20 (MaxPooling)	(None, 2, 2, 64)	0
dropout_14 (Dropout)	(None, 2, 2, 64)	0
flatten_6 (Flatten)	(None, 256)	0
dense_14 (Dense)	(None, 256)	65792
batch_normalization_7 (Batch)	(None, 256)	1024
dropout_15 (Dropout)	(None, 256)	0
dense_15 (Dense)	(None, 128)	32896
batch_normalization_8 (Batch)	(None, 128)	512
dropout_16 (Dropout)	(None, 128)	0

dense\_16 (Dense) (None, 10) 1290

---

Total params: 292,554  
Trainable params: 291,786  
Non-trainable params: 768

---

In [0]:

```
model.compile(optimizer='adam', loss='categorical_crossentropy', metrics=['accuracy'])
history = model.fit(x_train, y_train, batch_size=batch_size, epochs=epochs, verbose=1, validation_data=(x_test, y_test))
```

Train on 60000 samples, validate on 10000 samples

```
Epoch 1/12
60000/60000 [=====] - 672s 11ms/step - loss: 0.6168 - acc: 0.8232 - val_loss: 0.0631 - val_acc: 0.9816
Epoch 2/12
60000/60000 [=====] - 663s 11ms/step - loss: 0.0885 - acc: 0.9766 - val_loss: 0.0367 - val_acc: 0.9897
Epoch 3/12
60000/60000 [=====] - 662s 11ms/step - loss: 0.0605 - acc: 0.9844 - val_loss: 0.0385 - val_acc: 0.9898
Epoch 4/12
60000/60000 [=====] - 659s 11ms/step - loss: 0.0491 - acc: 0.9868 - val_loss: 0.0263 - val_acc: 0.9932
Epoch 5/12
60000/60000 [=====] - 659s 11ms/step - loss: 0.0422 - acc: 0.9891 - val_loss: 0.0232 - val_acc: 0.9937
Epoch 6/12
60000/60000 [=====] - 657s 11ms/step - loss: 0.0363 - acc: 0.9907 - val_loss: 0.0272 - val_acc: 0.9924
Epoch 7/12
60000/60000 [=====] - 647s 11ms/step - loss: 0.0309 - acc: 0.9916 - val_loss: 0.0264 - val_acc: 0.9928
Epoch 8/12
60000/60000 [=====] - 644s 11ms/step - loss: 0.0297 - acc: 0.9921 - val_loss: 0.0255 - val_acc: 0.9934
Epoch 9/12
60000/60000 [=====] - 646s 11ms/step - loss: 0.0237 - acc: 0.9934 - val_loss: 0.0312 - val_acc: 0.9909
Epoch 10/12
60000/60000 [=====] - 652s 11ms/step - loss: 0.0213 - acc: 0.9943 - val_loss: 0.0337 - val_acc: 0.9913
Epoch 11/12
60000/60000 [=====] - 660s 11ms/step - loss: 0.0210 - acc: 0.9946 - val_loss: 0.0250 - val_acc: 0.9934
Epoch 12/12
60000/60000 [=====] - 659s 11ms/step - loss: 0.0177 - acc: 0.9953 - val_loss: 0.0362 - val_acc: 0.9908
```


In [0]:

```
score = model.evaluate(x_test, y_test, verbose=0)
print('Test loss:', score[0])
print('Test accuracy:', score[1])
```

Test loss: 0.03619625380380894  
Test accuracy: 0.9908

In [0]:

```
fig, ax = plt.subplots(1, 1)
ax.set_xlabel('epoch') ; ax.set_ylabel('Categorical Crossentropy Loss')
x = list(range(1, epochs+1))
vy = history.history['val_loss']
ty = history.history['loss']
plt_dynamic(x, vy, ty, ax)
```



## 5 Convolution Layers

In [0]:

```
model = Sequential()

model.add(Conv2D(32, kernel_size=(3, 3),padding='same',activation='relu',input_shape=input_shape))
model.add(Conv2D(64,kernel_size = (5,5),padding = 'same',activation = 'relu'))
model.add(MaxPooling2D(pool_size=(3,3),padding='same'))

model.add(Conv2D(32, kernel_size=(4,4),padding='same',activation='relu'))
model.add(Conv2D(64,kernel_size = (5,5),padding='same', activation='relu'))
model.add(MaxPooling2D(pool_size=(3,3),padding='same'))

model.add(Conv2D(32, kernel_size=(5,5),padding='same',activation='relu'))
model.add(Conv2D(64,kernel_size = (6,6),padding='same', activation='relu'))
model.add(MaxPooling2D(pool_size=(3,3),padding='same'))

model.add(Conv2D(32, kernel_size=(4,4),padding='same',activation='relu'))
model.add(Conv2D(64,kernel_size = (6,6),padding='same', activation='relu'))
model.add(MaxPooling2D(pool_size=(3,3),padding='same'))

model.add(Conv2D(32, kernel_size=(5,5),padding='same',activation='relu'))
model.add(Conv2D(64,kernel_size = (6,6),padding='same', activation='relu'))
model.add(MaxPooling2D(pool_size=(3,3),padding='same'))

model.add(Dropout(0.5))

model.add(Flatten())

model.add(Dense(256,activation = 'relu'))
model.add(BatchNormalization())
model.add(Dropout(0.5))

model.add(Dense(256,activation = 'relu'))
model.add(BatchNormalization())
model.add(Dropout(0.5))

model.add(Dense(num_classes, activation='softmax'))

model.summary()
```

Layer (type)	Output Shape	Param #
conv2d_113 (Conv2D)	(None, 28, 28, 32)	320
conv2d_114 (Conv2D)	(None, 28, 28, 64)	51264
max_pooling2d_54 (MaxPooling)	(None, 10, 10, 64)	0
conv2d_115 (Conv2D)	(None, 10, 10, 32)	32800
conv2d_116 (Conv2D)	(None, 10, 10, 64)	51264
max_pooling2d_55 (MaxPooling)	(None, 4, 4, 64)	0

conv2d_117 (Conv2D)	(None, 4, 4, 32)	51232
conv2d_118 (Conv2D)	(None, 4, 4, 64)	73792
max_pooling2d_56 (MaxPooling)	(None, 2, 2, 64)	0
conv2d_119 (Conv2D)	(None, 2, 2, 32)	32800
conv2d_120 (Conv2D)	(None, 2, 2, 64)	73792
max_pooling2d_57 (MaxPooling)	(None, 1, 1, 64)	0
conv2d_121 (Conv2D)	(None, 1, 1, 32)	51232
conv2d_122 (Conv2D)	(None, 1, 1, 64)	73792
max_pooling2d_58 (MaxPooling)	(None, 1, 1, 64)	0
dropout_38 (Dropout)	(None, 1, 1, 64)	0
flatten_14 (Flatten)	(None, 64)	0
dense_38 (Dense)	(None, 256)	16640
batch_normalization_23 (Batch Normalization)	(None, 256)	1024
dropout_39 (Dropout)	(None, 256)	0
dense_39 (Dense)	(None, 256)	65792
batch_normalization_24 (Batch Normalization)	(None, 256)	1024
dropout_40 (Dropout)	(None, 256)	0
dense_40 (Dense)	(None, 10)	2570
=====		
Total params: 579,338		
Trainable params: 578,314		
Non-trainable params: 1,024		

In [0]:

```
model.compile(optimizer='adam', loss='categorical_crossentropy', metrics=['accuracy'])
history = model.fit(x_train, y_train, batch_size=batch_size, epochs=epochs, verbose=1, validation_data=(x_test, y_test))
```

Train on 60000 samples, validate on 10000 samples

```
Epoch 1/12
60000/60000 [=====] - 604s 10ms/step - loss: 1.4790 - acc: 0.3850 - val_loss: 1.0822 - val_acc: 0.5015
Epoch 2/12
60000/60000 [=====] - 602s 10ms/step - loss: 0.5329 - acc: 0.8220 - val_loss: 0.3030 - val_acc: 0.9029
Epoch 3/12
60000/60000 [=====] - 602s 10ms/step - loss: 0.2093 - acc: 0.9533 - val_loss: 0.1224 - val_acc: 0.9765
Epoch 4/12
60000/60000 [=====] - 604s 10ms/step - loss: 0.1478 - acc: 0.9695 - val_loss: 0.0933 - val_acc: 0.9807
Epoch 5/12
60000/60000 [=====] - 602s 10ms/step - loss: 0.1011 - acc: 0.9798 - val_loss: 0.0652 - val_acc: 0.9860
Epoch 6/12
60000/60000 [=====] - 602s 10ms/step - loss: 0.0941 - acc: 0.9813 - val_loss: 0.0956 - val_acc: 0.9787
Epoch 7/12
60000/60000 [=====] - 601s 10ms/step - loss: 0.0918 - acc: 0.9822 - val_loss: 0.0749 - val_acc: 0.9851
Epoch 8/12
60000/60000 [=====] - 600s 10ms/step - loss: 0.0851 - acc: 0.9825 - val_loss: 0.0588 - val_acc: 0.9884
Epoch 9/12
60000/60000 [=====] - 598s 10ms/step - loss: 0.0786 - acc: 0.9839 - val_loss: 0.0580 - val_acc: 0.9885
```

```

0.0000 - val_acc: 0.9900
Epoch 10/12
60000/60000 [=====] - 601s 10ms/step - loss: 0.0638 - acc: 0.9867 - val_loss:
0.0378 - val_acc: 0.9930
Epoch 11/12
60000/60000 [=====] - 610s 10ms/step - loss: 0.0480 - acc: 0.9907 - val_loss:
0.0488 - val_acc: 0.9893
Epoch 12/12
60000/60000 [=====] - 608s 10ms/step - loss: 0.0527 - acc: 0.9898 - val_loss:
0.0447 - val_acc: 0.9887

```

In [0]:

```

score = model.evaluate(x_test, y_test, verbose=0)
print('Test loss:', score[0])
print('Test accuracy:', score[1])

```

```

Test loss: 0.044668429520819335
Test accuracy: 0.9887


```

In [0]:

```

fig,ax = plt.subplots(1,1)
ax.set_xlabel('epoch') ; ax.set_ylabel('Categorical Crossentropy Loss')
x = list(range(1,epochs+1))
vy = history.history['val_loss']
ty = history.history['loss']
plt_dynamic(x, vy, ty, ax)

```



## 7 Convolutional Layers

In [0]:

```

model = Sequential()

model.add(Conv2D(32, kernel_size = (3,3), padding = 'same', activation = 'relu', input_shape = input_shape))
model.add(Conv2D(64, kernel_size = (5,5), padding = 'same', activation = 'relu'))
model.add(MaxPooling2D(pool_size=(3,3), padding='same'))

model.add(BatchNormalization())

model.add(Conv2D(32, kernel_size=(4,4), padding='same', activation='relu'))
model.add(Conv2D(64, kernel_size = (5,5), padding='same', activation='relu'))
model.add(MaxPooling2D(pool_size=(4,4), padding='same'))

model.add(BatchNormalization())

model.add(Conv2D(32, kernel_size=(5,5), padding='same', activation='relu'))
model.add(Conv2D(64, kernel_size = (6,6), padding='same', activation='relu'))
model.add(MaxPooling2D(pool_size=(5,5), padding='same'))

```

```

model.add(BatchNormalization())

model.add(Conv2D(32, kernel_size=(4,4),padding='same',activation='relu'))
model.add(Conv2D(64, kernel_size = (6,6),padding='same', activation='relu'))
model.add(MaxPooling2D(pool_size=(6,6),padding='same'))

model.add(BatchNormalization())

model.add(Conv2D(32, kernel_size=(5,5),padding='same',activation='relu'))
model.add(Conv2D(64, kernel_size = (6,6),padding='same', activation='relu'))
model.add(MaxPooling2D(pool_size=(3,3),padding='same'))

model.add(BatchNormalization())

model.add(Conv2D(32, kernel_size=(5,5),padding='same',activation='relu'))
model.add(Conv2D(64, kernel_size = (5,5),padding='same', activation='relu'))
model.add(MaxPooling2D(pool_size=(4,4),padding='same'))

model.add(BatchNormalization())

model.add(Conv2D(32, kernel_size=(4,4),padding='same',activation='relu'))
model.add(Conv2D(64, kernel_size = (7,7),padding='same', activation='relu'))
model.add(MaxPooling2D(pool_size=(5,5),padding='same'))

model.add(Dropout(0.5))

model.add(Flatten())

model.add(Dense(256,activation = 'relu'))
model.add(BatchNormalization())
model.add(Dropout(0.5))

model.add(Dense(256,activation = 'relu'))
model.add(BatchNormalization())
model.add(Dropout(0.5))

model.add(Dense(num_classes, activation='softmax'))

model.summary()

```

Layer (type)	Output Shape	Param #
conv2d_1 (Conv2D)	(None, 28, 28, 32)	320
conv2d_2 (Conv2D)	(None, 28, 28, 64)	51264
max_pooling2d_1 (MaxPooling2D)	(None, 10, 10, 64)	0
batch_normalization_1 (Batch Normalization)	(None, 10, 10, 64)	256
conv2d_3 (Conv2D)	(None, 10, 10, 32)	32800
conv2d_4 (Conv2D)	(None, 10, 10, 64)	51264
max_pooling2d_2 (MaxPooling2D)	(None, 3, 3, 64)	0
batch_normalization_2 (Batch Normalization)	(None, 3, 3, 64)	256
conv2d_5 (Conv2D)	(None, 3, 3, 32)	51232
conv2d_6 (Conv2D)	(None, 3, 3, 64)	73792
max_pooling2d_3 (MaxPooling2D)	(None, 1, 1, 64)	0
batch_normalization_3 (Batch Normalization)	(None, 1, 1, 64)	256
conv2d_7 (Conv2D)	(None, 1, 1, 32)	32800
conv2d_8 (Conv2D)	(None, 1, 1, 64)	73792
max_pooling2d_4 (MaxPooling2D)	(None, 1, 1, 64)	0
batch_normalization_4 (Batch Normalization)	(None, 1, 1, 64)	256
conv2d_9 (Conv2D)	(None, 1, 1, 32)	51232

conv2d_9 (Conv2D)	(None, 1, 1, 64)	73792
conv2d_10 (Conv2D)	(None, 1, 1, 64)	73792
max_pooling2d_5 (MaxPooling2D)	(None, 1, 1, 64)	0
batch_normalization_5 (Batch Normalization)	(None, 1, 1, 64)	256
conv2d_11 (Conv2D)	(None, 1, 1, 32)	51232
conv2d_12 (Conv2D)	(None, 1, 1, 64)	51264
max_pooling2d_6 (MaxPooling2D)	(None, 1, 1, 64)	0
batch_normalization_6 (Batch Normalization)	(None, 1, 1, 64)	256
conv2d_13 (Conv2D)	(None, 1, 1, 32)	32800
conv2d_14 (Conv2D)	(None, 1, 1, 64)	100416
max_pooling2d_7 (MaxPooling2D)	(None, 1, 1, 64)	0
dropout_1 (Dropout)	(None, 1, 1, 64)	0
flatten_1 (Flatten)	(None, 64)	0
dense_1 (Dense)	(None, 256)	16640
batch_normalization_7 (Batch Normalization)	(None, 256)	1024
dropout_2 (Dropout)	(None, 256)	0
dense_2 (Dense)	(None, 256)	65792
batch_normalization_8 (Batch Normalization)	(None, 256)	1024
dropout_3 (Dropout)	(None, 256)	0
dense_3 (Dense)	(None, 10)	2570
=====		
Total params: 816,586		
Trainable params: 814,794		
Non-trainable params: 1,792		
=====		

In [0]:

```
model.compile(optimizer='adam', loss='categorical_crossentropy', metrics=['accuracy'])
history = model.fit(x_train, y_train, batch_size=batch_size, epochs=epochs, verbose=1, validation_data=(x_test, y_test))
```

Train on 60000 samples, validate on 10000 samples

```
Epoch 1/12
60000/60000 [=====] - 603s 10ms/step - loss: 0.9687 - acc: 0.6452 - val_loss: 0.5443 - val_acc: 0.8350
Epoch 2/12
60000/60000 [=====] - 607s 10ms/step - loss: 0.2516 - acc: 0.9378 - val_loss: 0.1349 - val_acc: 0.9707
Epoch 3/12
60000/60000 [=====] - 621s 10ms/step - loss: 0.1385 - acc: 0.9714 - val_loss: 0.1593 - val_acc: 0.9672
Epoch 4/12
60000/60000 [=====] - 627s 10ms/step - loss: 0.1065 - acc: 0.9781 - val_loss: 0.0760 - val_acc: 0.9816
Epoch 5/12
60000/60000 [=====] - 580s 10ms/step - loss: 0.0860 - acc: 0.9825 - val_loss: 0.0876 - val_acc: 0.9794
Epoch 6/12
60000/60000 [=====] - 577s 10ms/step - loss: 0.0791 - acc: 0.9843 - val_loss: 0.1046 - val_acc: 0.9739
Epoch 7/12
60000/60000 [=====] - 578s 10ms/step - loss: 0.0708 - acc: 0.9856 - val_loss: 0.0503 - val_acc: 0.9885
Epoch 8/12
60000/60000 [=====] - 592s 10ms/step - loss: 0.0661 - acc: 0.9865 - val_loss: 0.0513 - val_acc: 0.9892
```



```
Epoch 9/12
60000/60000 [=====] - 597s 10ms/step - loss: 0.0619 - acc: 0.9874 - val_loss:
0.0408 - val_acc: 0.9912
Epoch 10/12
60000/60000 [=====] - 595s 10ms/step - loss: 0.0608 - acc: 0.9879 - val_loss:
0.0651 - val_acc: 0.9890
Epoch 11/12
60000/60000 [=====] - 612s 10ms/step - loss: 0.0512 - acc: 0.9897 - val_loss:
0.0411 - val_acc: 0.9922
Epoch 12/12
60000/60000 [=====] - 638s 11ms/step - loss: 0.0481 - acc: 0.9900 - val_loss:
0.0510 - val_acc: 0.9905
```

In [0]:

```
score = model.evaluate(x_test, y_test, verbose=0)
print('Test loss:', score[0])
print('Test accuracy:', score[1])
```

```
Test loss: 0.05102709092050791
Test accuracy: 0.9905
```

In [0]:

```
fig,ax = plt.subplots(1,1)
ax.set_xlabel('epoch') ; ax.set_ylabel('Categorical Crossentropy Loss')
x = list(range(1,epochs+1))
vy = history.history['val_loss']
ty = history.history['loss']
plt_dynamic(x, vy, ty, ax)
```

